

---

# ECE-GY 6143: Introduction to Machine Learning

## Midterm , Spring 2019

Prof. Sundeep Rangan

1. (13 points) *Least squares with a transformation.* Consider the following model for a scalar target  $\hat{y}$  from features  $\mathbf{x} = (x_1, x_2)$ ,

$$\hat{y} = f(x, \theta) = \begin{cases} \theta_0 + \theta_1(x_1 + \theta_2 x_2) & \text{if } x_1 < x_2 \\ \theta_0 + \theta_1(1 + \theta_2)x_1 & \text{if } x_1 \geq x_2 \end{cases}$$

where  $\theta = (\theta_0, \theta_1, \theta_2)$  are parameters.

- (a) (9 points) Write this as a linear model. That is, find basis functions  $\phi_i(\mathbf{x})$  and parameters  $\beta_i$  such that

$$\hat{y} = \sum_{i=1}^p \beta_i \phi_i(\mathbf{x}).$$

Use a minimum number  $p$  of basis functions. Write the linear model parameters  $\beta_i$  in terms of  $\theta_j$ .

- (b) (4 points) Given a linear parameter estimate  $\beta$ , how do you find  $\theta$  in the original model?

### Solution

- (a) Define,

$$\phi_1(\mathbf{x}) = 1, \quad \phi_2(\mathbf{x}) = x_1, \quad \phi_3(\mathbf{x}) = \begin{cases} x_2 & \text{if } x_1 < x_2 \\ x_1 & \text{if } x_1 \geq x_2. \end{cases}$$

Then,

$$f(x, \theta) = \theta_0 \phi_1(\mathbf{x}) + \theta_1 \phi_2(\mathbf{x}) + \theta_1 \theta_2 \phi_3(\mathbf{x}).$$

Define,

$$\beta_1 = \theta_0, \quad \beta_2 = \theta_1, \quad \beta_3 = \theta_1 \theta_2.$$

- (b) We invert the equations,

$$\theta_0 = \beta_1, \quad \theta_1 = \beta_2, \quad \theta_2 = \beta_3 / \theta_1 = \beta_3 / \beta_2.$$

2. (15 points) *Computing the bias.* Suppose that a true function is

$$y = f_0(x) := \min\{2x, 6\}.$$

We get data  $(x_i, y_i)$ ,  $i = 1, \dots, n$  with  $y_i = f_0(x_i)$ , and fit a model,

$$\hat{y} = f(\mathbf{x}, \hat{\theta}) = \hat{\theta}x, \quad \hat{\theta} = \frac{1}{n} \sum_{i=1}^n \frac{y_i}{x_i}.$$

For each set of the following sets of training data points  $\mathbf{x} = (x_1, \dots, x_n)$  and test data point  $x_{\text{ts}}$  find the bias:

$$\text{Bias}(x_{\text{ts}}) = f(x_{\text{ts}}, \hat{\theta}) - f_0(x_{\text{ts}}).$$

- (a) (5 points) Training data is  $\mathbf{x} = \{1, 2\}$  and test data point is  $x_{\text{ts}} = 2$ .
- (b) (5 points) Training data is  $\mathbf{x} = \{1, 2\}$  and test data point is  $x_{\text{ts}} = 4$ .
- (c) (5 points) Training data is  $\mathbf{x} = \{1, 4\}$  and test data point is  $x_{\text{ts}} = 4$ .

### Solution

For the three possible training points  $x_i = 1, 2, 4$ , we first compute  $y_i$  and  $y_i/x_i$  as shown in the table.

$x_i$	$y_i = \min\{2x, 6\}$	$y_i/x_i$
1	2	2
2	4	2
4	6	1.5

- (a) The parameter estimate  $\hat{\theta}$  is the average of  $y_i/x_i$  for the points in the training data. So, for the training data,  $\mathbf{x} = \{1, 2\}$ ,

$$\hat{\theta} = \frac{1}{2}(2 + 2) = 2,$$

and the bias at  $x_{\text{ts}} = 2$  is

$$\begin{aligned} \text{Bias}(x_{\text{ts}}) &= f(x_{\text{ts}}, \hat{\theta}) - f_0(x_{\text{ts}}) \\ &= \hat{\theta}x_{\text{ts}} - \max\{2x_{\text{ts}}, 6\} = 2(2) - 2(2) = 0. \end{aligned}$$

- (b) The training data is the same, so again we get  $\hat{\theta} = 2$ . The bias at  $x_{\text{ts}} = 4$  is

$$\begin{aligned} \text{Bias}(x_{\text{ts}}) &= f(x_{\text{ts}}, \hat{\theta}) - f_0(x_{\text{ts}}) \\ &= \hat{\theta}x_{\text{ts}} - \max\{2x_{\text{ts}}, 6\} = 4(2) - 6 = 2. \end{aligned}$$

- (c) For the training data,  $\mathbf{x} = \{1, 4\}$ ,

$$\hat{\theta} = \frac{1}{2}(2 + 1.5) = 1.75.$$

The bias at  $x_{\text{ts}} = 2$  is

$$\begin{aligned} \text{Bias}(x_{\text{ts}}) &= f(x_{\text{ts}}, \hat{\theta}) - f_0(x_{\text{ts}}) \\ &= \hat{\theta}x_{\text{ts}} - \max\{2x_{\text{ts}}, 6\} = 4(1.75) - 6 = 1. \end{aligned}$$

3. (15 points) *Model selection.* You are given python arrays  $\mathbf{X}$  and  $\mathbf{y}$  where  $\mathbf{X}$  is an  $n \times p$  data matrix and  $\mathbf{y}$  is a  $n$  dimensional vector of continuous targets. For each model order  $d=0,1,\dots,p$ , you consider a linear regression model using only the first  $d$  features:

$$\hat{y}[i] = b + \mathbf{X}[i,0]*w[0] + \dots + \mathbf{X}[i,d-1]*w[d-1],$$

where  $b$  and  $w[0],\dots,w[d-1]$  are the parameters. Write python code to:

- Split the data into training and test with approximately 50% of the samples in each.
- Use simple cross-validation to find the optimal order  $d$ .

You do not need to use  $K$ -fold validation or shuffle the data. You can use the normal rule, not the one SE rule. You may use the functions:

```
regr = LinearRegression() # Constructs a linear regression object
regr.fit(Z,y)             # Fits a model with features Z and outputs y
yhat = regr.predict(Z)    # Predicts outputs given features Z
```

**Solution.** One solution is as follows:

```
# Split the data into training and test
n = X.shape[0]
ntr = n // 2
Xtr = X[:ntr]
Xts = X[ntr:]
ytr = y[:ntr]
yts = y[ntr:]

# Loop over model orders
p = X.shape[1]
rss = np.zeros(p)
for d in range(p):
    # Fit linear model with d features
    regr = LinearRegression()
    regr.fit(Xtr[:, :d], ytr)

    # Predict on test
    yhat = regr.predict(Xts[:, :d])
    rss[d] = np.mean((yhat-yts)**2)

# Find optimal model order
dopt = np.argmin(rss)
```

4. (14 points) *Linear Classification.* You are given the following five training data points,  $(x_i, y_i)$ , with binary class labels  $y_i \in \{0, 1\}$ :

$x_i$	0	1.5	1.6	3	4
$y_i$	0	1	1	0	0

Consider a classifier of the form,

$$\hat{y}_i = \begin{cases} 1 & \text{if } z_i > 0 \\ 0 & \text{if } z_i < 0, \end{cases} \quad (1)$$

where  $z_i$  is some function of  $x_i$ .

- (a) (6 points) Find  $\beta_0, \beta_1$  such that with  $z_i = \beta_0 + \beta_1 x_i$ , the classifier (1) makes a minimum number of errors. Indicate which training data points, if any, are misclassified.
- (b) (8 points) Find  $\beta_0, \beta_1, \beta_2$  such that when  $z_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2$ , the classifier (1) makes a minimum number of errors. Indicate which training data points, if any, are misclassified.

### Solution

- (a) In this case, we are limited to a linear classifier. One possible linear classifier takes,

$$\hat{y} = 1 \text{ when } x < 2. \quad (2)$$

This makes only one error on the training data – the point  $(x_i, y_i) = (0, 0)$ . To implement the classifier (3), take

$$z_i = 2 - x_i,$$

which corresponds to  $(\beta_0, \beta_1) = (2, -1)$ .

- (b) To make no errors, we would like

$$\hat{y} = 1 \text{ when } x \in [1, 2]. \quad (3)$$

So, we would like  $z > 0$  in the region  $x \in (1, 2)$ . This can be done with the quadratic,

$$z = -(x - 1)(x - 2) = -x^2 + 3x - 2.$$

So, we take  $(\beta_0, \beta_1, \beta_2) = (-2, 3, -1)$ .

- 5. (14 points) *Regularized least squares*. We are given data  $(x_i, y_i)$ ,  $i = 0, \dots, T - 1$  and wish to fit a model of the form,

$$y_i \approx \hat{y}_i = \sum_{j=0}^{d-1} w_j x_{i-j},$$

for some parameters  $w_j$ . You can assume  $x_j = 0$  for  $j < 0$ . You decide to use regularized least squares,

$$\hat{\mathbf{w}} := \arg \min_{\mathbf{w}} \frac{1}{2} \|\mathbf{y} - \mathbf{A}\mathbf{w}\|^2 + \lambda \phi(\mathbf{w}).$$

- (a) (5 points) What are the components  $A_{ij}$  of the matrix  $\mathbf{A}$ ?
- (b) (5 points) Suggest a function  $\phi(\mathbf{w})$  if it is known that we should have  $w_j = 0$  for most  $j$ .
- (c) (4 points) Suggest a function  $\phi(\mathbf{w})$  if it is known that we should have  $w_j \geq 0$  and  $w_{j+1} \leq w_j$  for most indices  $j$ .

For parts (b) and (c), there is no single correct answer.

### Solution

- (a) Components are  $A_{ij} = x_{i-j}$  with

$$\mathbf{A} = \begin{bmatrix} x_0 & 0 & \cdots & 0 \\ x_1 & x_0 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ x_{T-1} & x_{T-2} & \cdots & x_{T-d} \end{bmatrix}.$$

To get full credit, you needed to state the dimensions of  $\mathbf{A} \in \mathbb{R}^{T \times d}$ .

- (b) Take  $\phi(\mathbf{w}) = \|\mathbf{w}\|_1$ .  
 (c) There are a few ways you can do this. For example, you can take a function such as

$$\phi(\mathbf{w}) := \sum_{j=0}^{d-1} \max\{0, -w_j\} + \sum_{j=0}^{d-2} \max\{0, w_{j+1} - w_j\}.$$

The first term penalizes negative values of  $w_j$ . The second term penalizes positive values  $w_{j+1} - w_j$ .

6. (16 points) *Computing gradients.* Consider the function,

$$J = \sum_{i=1}^N \ln(1 + e^{z_i}) - z_i y_i, \quad z_i = \sum_{j=1}^p \frac{a_j}{1 + (x_i - b_j)^2}.$$

- (a) (8 points) Compute the gradient components,  $\partial J / \partial a_j$  and  $\partial J / \partial b_j$ .  
 (b) (8 points) Write a python function to that returns function  $J$  and the gradients,

```
def Jval(...):
    ...
    return J, Jgrad_a, Jgrad_b
```

You must determine the arguments for the function. For full credit, avoid for loops.

### Solution

- (a) Let  $L_i = \ln(1 + e^{z_i}) - z_i y_i$  so that  $L = \sum_i L_i$  and

$$\frac{\partial L_i}{\partial z_i} = \frac{e^{z_i}}{1 + e^{z_i}} - y_i = \frac{1}{1 + e^{-z_i}} - y_i.$$

Then,

$$\frac{\partial L}{\partial a_j} = \sum_{i=1}^N \frac{\partial L_i}{\partial a_j} = \sum_{i=1}^N \frac{\partial L_i}{\partial z_i} \frac{\partial z_i}{\partial a_j} = \sum_{i=1}^N \frac{\partial L_i}{\partial z_i} \frac{1}{1 + (x_i - b_j)^2}.$$

Similarly,

$$\frac{\partial L}{\partial b_j} = \sum_{i=1}^N \frac{\partial L_i}{\partial b_j} = \sum_{i=1}^N \frac{\partial L_i}{\partial z_i} \frac{\partial z_i}{\partial b_j} = \sum_{i=1}^N \frac{\partial L_i}{\partial z_i} \frac{2a_j(x_i - b_j)}{(1 + (x_i - b_j)^2)^2}.$$

(b) One possible solution is as follows:

```
def feval(a,b,x,y):
    # Compute difference D[i,j] = x[i] - b[j]
    D = x[:,None] - b[None,:]

    # Compute P[i,j] = 1/(1+D[i,j]**2)
    #           z[i] = sum_j P[i,j]*a[j]
    P = 1/(1+D**2)
    z = np.sum(a[None,:]*P,axis=1)

    # Compute loss
    J = np.sum(np.log(1+np.exp(z)) - z*y)

    # Compute dJ_dz[i]
    dJ_dz = 1/(1+np.exp(-z)) - y

    # Compute gradients
    Jgrad_a = np.sum(dJ_dz[None,:]*P,axis=0)
    Jgrad_b = np.sum(2*dJ_dz[None,:]*a[None,:]*D*(P**2),axis=0)

    return J, Jgrad_a, Jgrad_b
```

7. (13 points) *Gradient descent with a stopping condition.* You are given a python function that returns a loss function  $f(\mathbf{a}, \mathbf{b})$  and the gradients with respect to vector parameters  $\mathbf{a}$  and  $\mathbf{b}$ :

```
f, fgrad_a, fgrad_b = feval(a,b)
```

Write a few lines of python code that runs gradient descent for some initial conditions `ainit` and `binit` with a fixed step size `step`. The code is suppose to stop when either of the following conditions occurs:

- `max_iter=500` iterations have been done, or
- $\sum_j (\partial f / \partial a_j)^2 \leq (10)^{-8}$  and  $\sum_k (\partial f / \partial b_k)^2 \leq (10)^{-8}$ .

In python, the format for a `while` loop is:

```
while not done:
    # do something
```

**Solution.** One solution is as follows:

```
a = ainit
b = binit
tol = 1e-8
step = ...
max_iter = 500
it = 0
done = False
```

```
while not done:
    f, fgrad_a, fgrad_b = feval(a,b)
    a = a - step*fgrad_a
    b = b - step*fgrad_b

    # Check stopping condition
    it += 1
    norma = np.sum(fgrad_a**2)
    normb = np.sum(fgrad_b**2)
    done = (it >= max_iter) or ((norma < tol) and (normb < tol))
```